

Lab2MD

April 9, 2025

```
[3]: # Data manipulation and visualization
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Model development
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer

# Models
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree

# Model evaluation
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    classification_report,
    confusion_matrix
)

# For displaying results nicely in the notebook
from IPython.display import display, HTML

# Set visualization style
plt.style.use('seaborn-v0_8-whitegrid')
sns.set_palette('viridis')
```

```
# For reproducibility
import random
random.seed(42)
np.random.seed(42)
```

```
[4]: # Read the student performance dataset
df = pd.read_csv('student_performance_data.csv')

# Remove StudentID column from the dataset, irrelevant for what I am finding
df = df.drop(columns=['StudentID'])

# Basic info about the dataset
print("Dataset shape:", df.shape)
print("\nFirst 5 rows of the dataset:")
display(df.head())

# Descriptive statistics
print("\nDescriptive statistics:")
display(df.describe())

# Check for missing values
print("\nMissing values per column:")
display(df.isnull().sum())

# Display information about data types and non-null counts
print("\nDataset information:")
display(df.info())
```

Dataset shape: (500, 8)

First 5 rows of the dataset:

| | Gender | Age | StudyHoursPerWeek | AttendanceRate | GPA | Major \ |
|---|--------|-----|-------------------|----------------|------|-----------|
| 0 | Male | 24 | 37 | 90.75 | 3.47 | Arts |
| 1 | Female | 22 | 37 | 74.90 | 2.32 | Education |
| 2 | Male | 22 | 10 | 53.36 | 2.38 | Business |
| 3 | Male | 24 | 10 | 70.26 | 3.46 | Science |
| 4 | Male | 18 | 19 | 74.87 | 2.31 | Education |

| | PartTimeJob | ExtraCurricularActivities |
|---|-------------|---------------------------|
| 0 | Yes | No |
| 1 | No | No |
| 2 | No | No |
| 3 | Yes | No |
| 4 | Yes | No |

Descriptive statistics:

| | Age | StudyHoursPerWeek | AttendanceRate | GPA |
|-------|------------|-------------------|----------------|------------|
| count | 500.000000 | 500.000000 | 500.000000 | 500.000000 |
| mean | 20.956000 | 19.876000 | 74.990380 | 2.98516 |
| std | 2.000517 | 11.471347 | 14.565917 | 0.56362 |
| min | 18.000000 | 1.000000 | 50.010000 | 2.00000 |
| 25% | 19.000000 | 10.000000 | 62.607500 | 2.48750 |
| 50% | 21.000000 | 20.500000 | 75.730000 | 3.00000 |
| 75% | 23.000000 | 30.000000 | 87.220000 | 3.48000 |
| max | 24.000000 | 39.000000 | 99.970000 | 3.99000 |

Missing values per column:

```

Gender          0
Age             0
StudyHoursPerWeek  0
AttendanceRate  0
GPA             0
Major           0
PartTimeJob     0
ExtraCurricularActivities  0
dtype: int64

```

Dataset information:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 500 entries, 0 to 499

Data columns (total 8 columns):

| # | Column | Non-Null Count | Dtype |
|---|---------------------------|----------------|---------|
| 0 | Gender | 500 non-null | object |
| 1 | Age | 500 non-null | int64 |
| 2 | StudyHoursPerWeek | 500 non-null | int64 |
| 3 | AttendanceRate | 500 non-null | float64 |
| 4 | GPA | 500 non-null | float64 |
| 5 | Major | 500 non-null | object |
| 6 | PartTimeJob | 500 non-null | object |
| 7 | ExtraCurricularActivities | 500 non-null | object |

dtypes: float64(2), int64(2), object(4)

memory usage: 31.4+ KB

None

```

[5]: # Step 1: Create binary target based on GPA
gpa_column = 'GPA'
df['target'] = (df[gpa_column] >= 3.2).astype(int)

# Display the distribution of the target variable
print("Target distribution:")

```

```

print(df['target'].value_counts())
print(f"Percentage of students with GPA 3.2: {df['target'].mean()*100:.2f}%")

# Step 2: Prepare the data
# Exclude GPA and target from features
X = df.drop(columns=[gpa_column, 'target'])
y = df['target']

# Handle categorical features
categorical_features = X.select_dtypes(include=['object', 'category']).columns.
    ↪tolist()
numerical_features = X.select_dtypes(include=['int64', 'float64']).columns.
    ↪tolist()

# Create preprocessing pipelines
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

# Combine preprocessing steps
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 3: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42, stratify=y)

# Step 4: Build the models
# Logistic Regression
log_reg_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(max_iter=1000, random_state=42))
])

# k-NN
knn_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', KNeighborsClassifier(n_neighbors=5))
])

```

```

])

# Decision Tree
dt_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', DecisionTreeClassifier(random_state=42))
])

# Dictionary of models
models = {
    'Logistic Regression': log_reg_pipeline,
    'k-Nearest Neighbors': knn_pipeline,
    'Decision Tree': dt_pipeline
}

# Step 5: Train and evaluate each model
results = {}

for name, model in models.items():
    print(f"\n{'-'*50}")
    print(f"Training and evaluating {name}...")

    # Train the model
    model.fit(X_train, y_train)

    # Make predictions
    y_pred = model.predict(X_test)

    # Calculate metrics
    accuracy = accuracy_score(y_test, y_pred)

    # Store results
    results[name] = {
        'accuracy': accuracy,
        'predictions': y_pred
    }

    # Print classification report
    print(f"\nClassification Report for {name}:")
    print(classification_report(y_test, y_pred))

# Create and display confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['GPA < 3.2', 'GPA >= 3.2'],
            yticklabels=['GPA < 3.2', 'GPA >= 3.2'])

```

```

plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix - {name}')
plt.tight_layout()
plt.show()

# Step 6: Compare model performance
model_comparison = pd.DataFrame({
    'Model': list(results.keys()),
    'Accuracy': [results[model]['accuracy'] for model in results]
})

# Sort by accuracy
model_comparison = model_comparison.sort_values('Accuracy', ascending=False).
    ↪reset_index(drop=True)

# Display comparison
print("\nModel Comparison:")
display(model_comparison)

# Visualize model comparison
plt.figure(figsize=(10, 6))
sns.barplot(x='Model', y='Accuracy', data=model_comparison)
plt.title('Model Accuracy Comparison')
plt.ylim(0, 1)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

Target distribution:

target

0 300

1 200

Name: count, dtype: int64

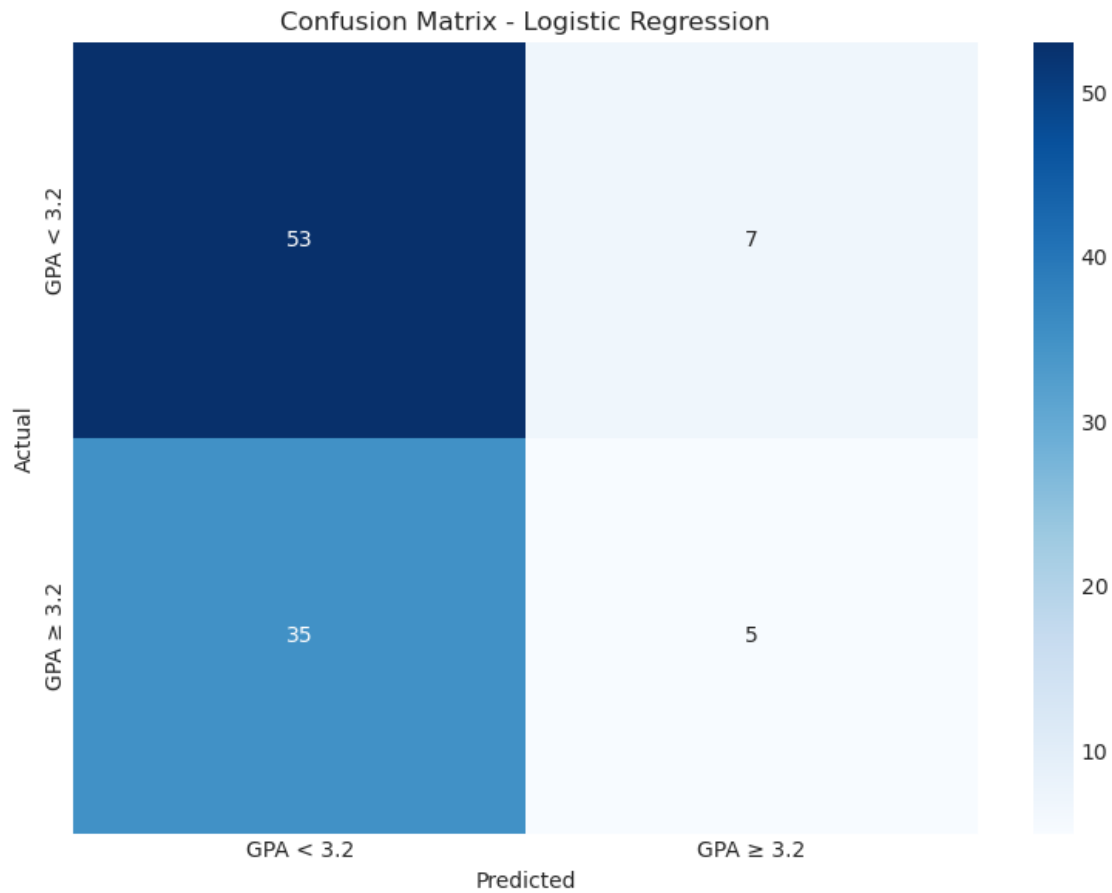
Percentage of students with GPA 3.2: 40.00%

Training and evaluating Logistic Regression...

Classification Report for Logistic Regression:

| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| 0 | 0.60 | 0.88 | 0.72 | 60 |
| 1 | 0.42 | 0.12 | 0.19 | 40 |
| accuracy | | | 0.58 | 100 |
| macro avg | 0.51 | 0.50 | 0.45 | 100 |

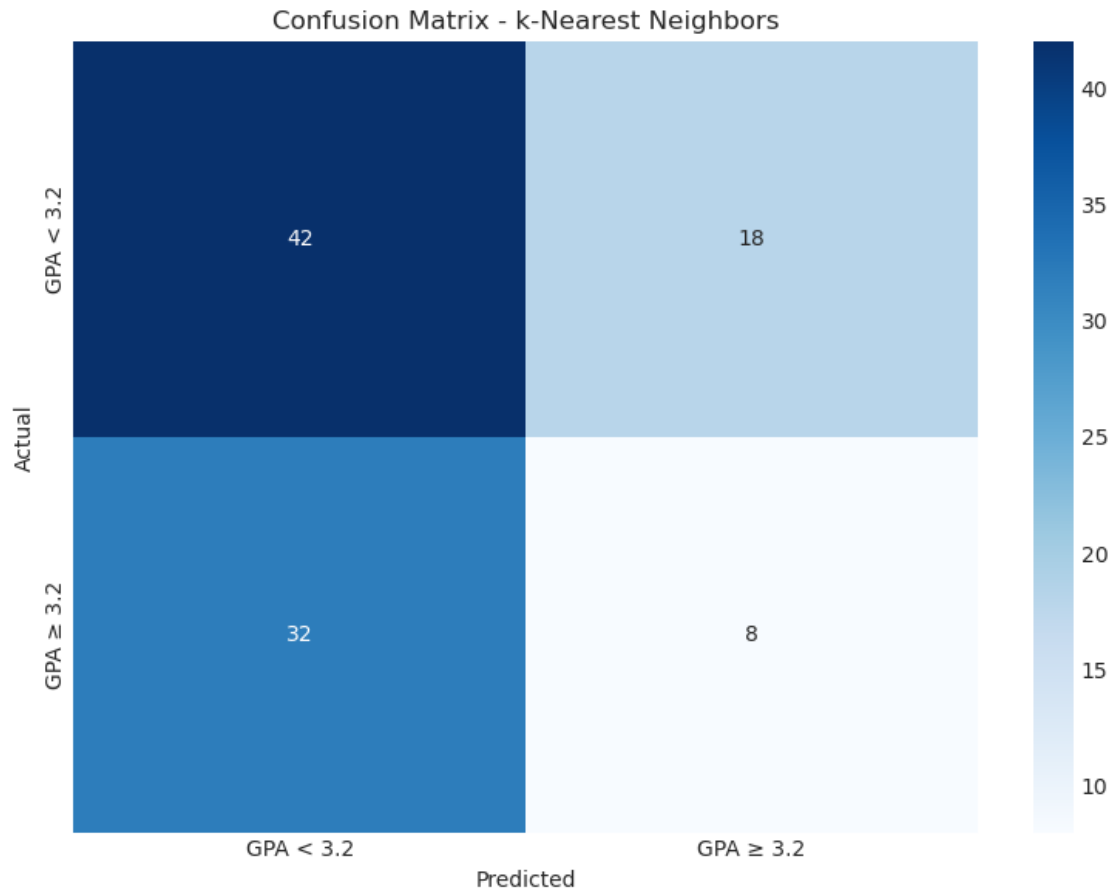
| | | | | |
|--------------|------|------|------|-----|
| weighted avg | 0.53 | 0.58 | 0.51 | 100 |
|--------------|------|------|------|-----|



 Training and evaluating k-Nearest Neighbors...

Classification Report for k-Nearest Neighbors:

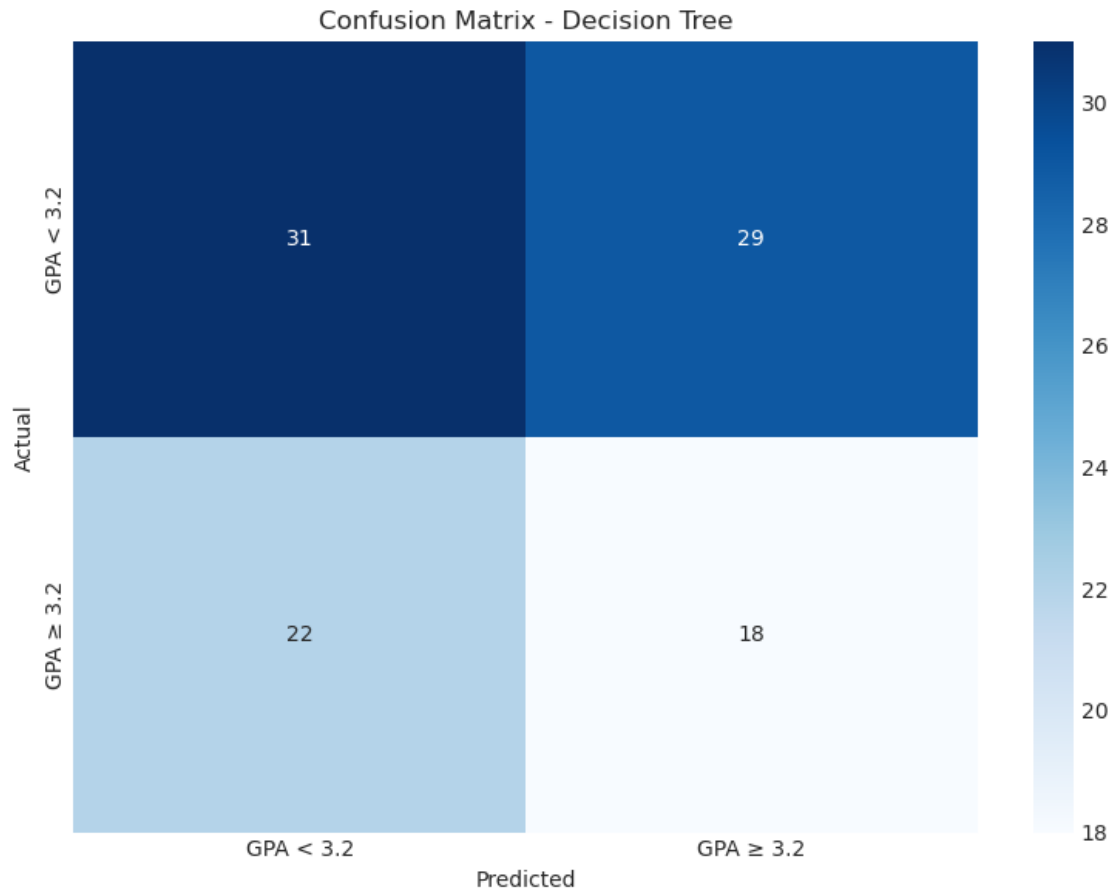
| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.57 | 0.70 | 0.63 | 60 |
| 1 | 0.31 | 0.20 | 0.24 | 40 |
| accuracy | | | 0.50 | 100 |
| macro avg | 0.44 | 0.45 | 0.43 | 100 |
| weighted avg | 0.46 | 0.50 | 0.47 | 100 |



Training and evaluating Decision Tree...

Classification Report for Decision Tree:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.58 | 0.52 | 0.55 | 60 |
| 1 | 0.38 | 0.45 | 0.41 | 40 |
| accuracy | | | 0.49 | 100 |
| macro avg | 0.48 | 0.48 | 0.48 | 100 |
| weighted avg | 0.50 | 0.49 | 0.49 | 100 |



Model Comparison:

| | Model | Accuracy |
|---|---------------------|----------|
| 0 | Logistic Regression | 0.58 |
| 1 | k-Nearest Neighbors | 0.50 |
| 2 | Decision Tree | 0.49 |

